



Automatic Hate Speech Detection on Social Media Text using Machine Learning

J J C Prasad Yadav | Chimmani Aruna | N Srilatha

Independent Researchers

Corresponding Author Mail Id: jagadishkrishnan0527@gmail.com, arunachimmani@gmail.com, srilathancse@gmail.com

To Cite this Article

J J C Prasad Yadav, Chimmani Aruna and N Srilatha. Automatic Hate Speech Detection on Social Media Text using Machine Learning. *International Journal for Modern Trends in Science and Technology* 2022, 8 pp. 124-129.

<https://doi.org/10.46501/IJMTST0801022>

Article Info

Received: 16 November 2021; Accepted: 01 January 2022; Published: 06 January 2022

ABSTRACT

The increasing use of social media and information sharing has given major benefits to humanity. However, this has also given rise to a variety of challenges including the spreading and sharing of hate speech messages. Hate speech is one type of harmful online content which directly attacks or promotes hate towards a group, or an individual member based on their actual or perceived aspects of identity, such as ethnicity, religion, and sexual orientation. With online hate speech on the rise, its automatic detection as a natural language processing task is gaining increasing interest. In our work, we included pre-processing (stop word removal and stemming) and multiple classifiers for performance evaluation. The results we obtained were satisfactory.

Keywords: social media, Hate speech, Abusive language, machine learning, natural language processing and Text classification.

1. INTRODUCTION

The debate around the regulation of hate speech is still ongoing [1]. It is still not clear whether the best response to it is through legal measures, or other methods (such as counter-speech and education [2]). Regardless of the means of countering it, the evident harm of hate speech [3] makes its detection crucial. Both the volume of content generated online, particularly in social media, and the psychological burden of manual moderation [4] supports the need for the automatic detection of offensive and hateful content. Online hate, described as abusive language [5], aggression [6], cyberbullying [7, 8], hatefulness [9], insults [10], personal attacks [11], provocation [12], racism [13], sexism [14], threats [15], or toxicity [16], has been identified as a major threat on online social media platforms. Pew Research Centre [17] reports that among 4248 adults in the United States, 41% have personally

experienced harassing behaviour online, whereas 66% witnessed harassment directed towards others. Around 22% of adults have experienced offensive name-calling, purposeful embarrassment (22%), physical threats (10%), and sexual harassment (6%), among other types of harassment. Social media platforms are the most prominent grounds for such toxic behaviour. Even though they often provide ways of flagging offensive and hateful content, only 17% of all adults have flagged harassing conversation, whereas only 12% of adults have reported someone for such acts [17]. In the following paragraphs, we better tried to differentiate Hate speech, offensive language, and abusive language. Although different types of abusive and offensive language are closely related, there are important distinctions to note. Offensive language and abusive language are both used as umbrella terms for

harmful content in the context of automatic detection studies. However, while “strongly impolite, rude” and possible use of profanity are seen in the definitions of both, abusive language has a strong component of intentionality. Thus, offensive language has a broader scope, and hate speech falls in both categories. Because of its definition mentioned above, hate speech is also different from other subtypes of offensive language. For example, personal attacks are characterised by being directed at an individual, which is not necessarily motivated by the target’s identity. Hate speech is also different from cyberbullying which is carried out repeatedly and over time against vulnerable victims that cannot defend themselves. This paper focuses on hate speech and hate speech datasets, although studies that cover both hate speech and other offensive language are also mentioned.

2.RELATED WORK

Abusive messages in social media is a complex phenomenon with a broad range of overlapping modes and goals [18]. Cyberbullying and hate speech are typical examples of abusive languages that researchers have put more interest in the past few decades due to their negative impacts in our societies. Several research have been conducted to automatically detect these undesirable messages among other messages in social media. The automatic detection of hate speech using machine learning approaches is relatively new, and there are very limited review papers on techniques for automatic hate speech detection [19].

The recent and related survey papers available on review of hate speech detection methods during this research work were few. The following were the available traditional literature review related to automatic detection of hate speech using ML: [20], ML algorithms have contributed immensely to hate speech detection and SM content analysis generally [21]. Offensive comments such as HS and cyberbullying are the most researched areas in NLP in the past few decades [22].

ML algorithms have been of great help in this direction in terms of SM data analysis for the identification and classification of offensive comments [23]. The advances in ML algorithms research have made significant impacts in many fields of endeavour which led to some important tools and models for analysing a large amount

of data in real-world problems like SMNs content analysis [24]. In this survey conducted by [25], the authors presented a brief review on eight hate speech detection techniques and approaches. These eight techniques include TF-IDF, dictionaries, N-gram, sentiment analyses, template-based approach, part of speech, Bag of the word, and rule-based approach. The limitation of the review is that techniques such as deep learning and ensemble approach were not considered in their work. In [19], the authors offered a brief, and critical analysis of the areas of automated hate speech detection in natural language processing. The authors also analysed the features for hate speech detection in literature which includes simple surface features, word generalization, sentiment analysis, lexical resources, linguistic features, knowledge-based features, meta-information, and multimodal information. The limitation of these two reviews is that techniques such as deep learning and ensemble approach are not considered in their work. The most significant step in text classification pipeline is selection of the best classifier [26]. Therefore, the need to review all techniques is of essence. We intend to make this selection phase easier for researchers by reviewing more algorithms than the previous review work has covered. In this case, we reviewed techniques like deep learning, ensemble learning among others that have been employed for the automatic detection of hate speech in social media. Posters of hate speeches usually attack their targets using the following attributes: Religion, Race, political affiliation, gender, marital status, ethnicity, health status, disability, and nationality [27]. The following diagram shows the worldwide most popular social networks as of October 2021[28], ranked by number of active users (in millions).




Fig.1: Most popular social networks worldwide (as of October 2021), ranked by number of active users (in millions)

3.Challenges of Detecting Hateful and Offensive Speech

There are many layers to the difficulty of automatically detecting hateful and/or offensive speech, particularly in social media. Some of these difficulties being closely related to the shortcomings of keyword-based approaches. For one, words can be obfuscated in many ways, both in an intentional attempt to avoid automatic content moderation, or because of the use of social media for communication (consider, for example the tendency in some posts to replace letters with similar looking numbers, e.g., “E”s with 3s, or “I”s with 1s, and so on). Furthermore, there are many expressions that are not inherently offensive, however they can be so in the right context. But even in the case of slurs, not only different slurs hold a different degree of offense, but the offense can also vary based on different time (as previously innocuous words may become slurs in time), as well as different use of the same word, different users, and different audience members. One example of this is the difference in the use of slurs by in-group speakers, and out-group speakers. This factor, when disregarded can contribute to the bias in hate speech detection corpora (against African Americans, and more specifically, against African American men), and in turn, the bias in hate speech detection (a strong argument for the transparency and explainability of hate speech detection models).

One recommendation to mitigate bias is explicitly preparing annotators for it. This leads to another difficulty, namely the availability (or lack thereof) of reliably annotated data. A factor that contributes to this problem is that there is no universally accepted definition of hate speech (a statement many publications would agree on, let alone one that is productive. One can point at a United Nations report for definition, we would however argue that it does not satisfy the criteria of being a universally accepted productive definition on several accounts. For one, the recommendations in said document are not legally binding, thus their implementation in all member countries is not a given. Furthermore, the recommendation here is only to “draw [...] from the guidance and definitions”, not to apply them as it is, thus even if the recommendations were binding, or all member countries would decide by their own volition to accept them, different countries

could still arrive at different definitions implemented in their domestic legal frameworks. Moreover, even if the definitions were used “as is”, the question remains whether they would be applicable for large scale data annotation, considering the contextual nature of their terms (“First, one should realize that the question of distinguishing those forms of expression that should be defined as incitement to hatred and thus prohibited is contextual and the individual circumstances and the individual circumstances of each case, such as local conditions, history, cultural and political tensions, must be considered”), and the complexity of definitions that could necessitate annotators having a background in law.

One benefit of a universally agreed upon productive definition for hate speech could be important for more reliable annotation, with higher inter-annotator agreement. For example, the 2019 HASOC hate speech and offensive content evaluation task had an interrater agreement rate that is between 69 and 77 percent for different task, even though “many texts recommend 80agreement” [30]. According to Mandl et al. [31] (the organizers of the 2019 HASOC hate speech and offensive content evaluation task), one difficulty in annotation (an issue that may have contributed to the low interrater agreements) was the use of language registers, such as youth talk. The difficulty of annotating youth talk is exemplified by the annotation of some example tweets (see Table 1) where the name of India’s prime minister (Narendra Modi, or Modi Ji) was used in various pop-cultural references (or “memes”). The first being a paraphrase of the chorus (“Never gonna give you up // Never gonna let you down // Never gonna run around and desert you..”) from the 1987 Rick Astley hit, Never Gonna Give You Up (that gained a considerable reputation in recent years, due to its use in the phenomenon called “rickrolling”). The second being a reference to a popular beverage that is well known for the company’s secrecy regarding its recipe. While the third referencing a much-quoted part of a recent movie, and the last one making a reference to a 1998 song from Baha Man (Who let the dogs out). Despite the similar nature of the tweets (particularly the last three tweets, as all three of them allude to Modi Ji knowing something that in general considered impenetrable - as mentioned before, the ingredients of Coca Cola are

- [6] Kumar S, et al. Community interaction and conflict on the web. In: Proceedings of the 2018 world wide web conference on world wide web; 2018. P. 933–43.
- [7] Hosseinmardi H et al (2015) Analyzing labeled cyberbullying incidents on the instagram social network. Soc Inf 2015:49–66. [8] Wachs S et al (2019) Understanding the overlap between cyberbullying and cyberhate perpetration: moderating effects of toxic online disinhibition. Crim Behav Mental Health 29(3):179–188.
- [9] Salminen J, et al. Anatomy of online hate: developing a taxonomy and machine learning models for identifying and classifying hate in online news media. In: Proceedings of the international AAAI conference on web and social media (ICWSM 2018), San Francisco; 2018.
- [10] Sood SO et al (2012) Automatic identification of personal insults on social news sites. J Am Soc Inform Sci Technol 63(2):270–285.
- [11] Wulczyn E, et al. Ex Machina: personal attacks seen at scale. In: Proceedings of the 26th international conference on world wide web, Geneva; 2017. P. 1391–9.
- [12] Mkono M (2018) ‘Troll alert!’: provocation and harassment in tourism and hospitality social media. Curr Issues Tour 21(7):791–804.
- [13] Waseem Z. Are you a racist or am i seeing things? Annotator influence on hate speech detection on twitter. In: Proceedings of the first workshop on NLP and computational social science; 2016. P. 138–42.
- [14] Chatzakou D, et al. Measuring #GamerGate: A tale of hate, sexism, and bullying. In: Proceedings of the 26th international conference on world wide web companion, Geneva; 2017. P. 1285–90.
- [15] Willard NE (2007) Cyberbullying and cyberthreats: Responding to the challenge of online social aggression, threats, and distress. research press, Champaign.
- [16] Märtens M, et al. Toxicity detection in multiplayer online games. In: Proceedings of the 2015 international workshop on network and systems support for games, Piscataway; 2015. P. 5:1–5:6.
- [17] Pew Research Center 2017. Online Harassment 2017.
- [18] R. Alshalan and H. Al-Khalifa, “A deep learning approach for automatic hate speech detection in the Saudi Twittersphere,” Appl. Sci., vol. 10, no. 23, pp. 1–16, 2020.
- [19] A. Al-Hassan and H. Al-Dossari, “Detection of hate speech in social networks: A survey on multilingual corpus,” in Proc. Comput. Sci. Inf. Technol. (CS IT), Feb. 2019, pp. 83–100.
- [20] A. Schmidt and M. Wiegand, “A survey on hate speech detection using natural language processing,” in Proc. 5th Int. Workshop Natural Lang. Process. Social Media, 2017, pp. 1–10.
- [21] M. A. Al-Garadi, M. R. Hussain, N. Khan, G. Murtaza, H. F. Nweke, I. Ali, G. Mujtaba, H. Chiroma, H. A. Khattak, and A. Gani, “Predicting cyberbullying on social media in the big data era using machine learning algorithms: Review of literature and open challenges,” IEEE Access, vol. 7, pp. 70701–70718, 2019.
- [22] Rodriguez, C. Argueta, and Y.-L. Chen, “Automatic detection of hate speech on Facebook using sentiment and emotion analysis,” in Proc. Int. Conf. Artif. Intell. Inf. Commun. (ICAIC), Feb. 2019, pp. 169–174.
- [23] G. Weir, K. Owuoye, A. Oberacker, and H. Alshahrani, “Cloud-based textual analysis as a basis for document classification,” in Proc. Int. Conf. High Perform. Comput. Simul. (HPCS), Jul. 2018, pp. 629–633.
- [24] J. Cheng, C. Danescu-Niculescu-Mizil, and J. Leskovec, “Anti-social behavior in online discussion communities,” in Proc. 9th Int. Conf. Web Soc. Media (ICWSM), 2015, pp. 61–70, 2015.
- [25] A. Alrehili, “Automatic hate speech detection on social media: A brief survey,” in Proc. IEEE/ACS 16th Int. Conf. Comput. Syst. Appl. (AICCSA), Nov. 2019, pp. 1–6.
- [26] K. Kowsari, K. J. Meimandi, M. Heidarysafa, S. Mendu, L. Barnes, and D. Brown, “Text classification algorithms: A survey,” Information, vol. 10, no. 4, pp. 1–68, 2019.
- [27] T. Granskogen and J. A. Gulla, “Fake news detection: Network data from social media used to predict fakes,” in Proc. CEUR Workshop, vol. 2041, no. 1, 2017, pp. 59–66.
- [28] <https://www.statista.com/statistics/272014/global-social-networks-ranked-by-number-of-users/>
- [29] Davidson et al., “Automated Hate Speech Detection and the Problem of Offensive Language”, Proceedings of the 11th International AAAI Conference on Web and Social Media, 2017.
- [30] McHugh M. Interrater reliability: the kappa statistic. Biochemia medica : časopis Hrvatskoga društva medicinskih biokemičara / HDMB. 2012;22:276–82.
- [31] Mandl T, Modha S, Patel D, Dave M, Mandlia C, Patel A. Overview of the HASOC track at FIRE 2019: Hate Speech and Offensive Content Identification in Indo-European Languages). In: Proceedings of the 11th annual meeting of the Forum for Information Retrieval Evaluation; 2019.